

Multiclass Support Vector Machine exercise

In this exercise you will:

- implement a fully-vectorized **loss function** for the SVM
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** using numerical gradient
- use a validation set to **tune the learning rate and regularization** strength
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

In [1]:

```
1 # Run some setup code for this notebook.
2
3 import random
4 import numpy as np
5 from cs175.data_utils import load_CIFAR10
6 import matplotlib.pyplot as plt
7
8
9 # This is a bit of magic to make matplotlib figures appear inline in the
10 # notebook rather than in a new window.
11 %matplotlib inline
12 plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
13 plt.rcParams['image.interpolation'] = 'nearest'
14 plt.rcParams['image.cmap'] = 'gray'
15
16 # Some more magic so that the notebook will reload external python modules;
17 # see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
18 %load_ext autoreload
19 %autoreload 2
```

CIFAR-10 Data Loading and Preprocessing

In [2]:

```
1 # Load the raw CIFAR-10 data.
2 cifar10_dir = 'cs175/datasets/cifar-10-batches-py'
3 X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)
4
5 # As a sanity check, we print out the size of the training and test data.
6 print('Training data shape: ', X_train.shape)
7 print('Training labels shape: ', y_train.shape)
8 print('Test data shape: ', X_test.shape)
9 print('Test labels shape: ', y_test.shape)
```

Training data shape: (50000, 32, 32, 3)


Training labels shape: (50000,)

Test data shape: (10000, 32, 32, 3)

Test labels shape: (10000,)

In [3]:

```
1 # Visualize some examples from the dataset.
2 # We show a few examples of training images from each class.
3 classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
4 num_classes = len(classes)
5 samples_per_class = 7
6 for y, cls in enumerate(classes):
7     idxs = np.flatnonzero(y_train == y)
8     idxs = np.random.choice(idxs, samples_per_class, replace=False)
9     for i, idx in enumerate(idxs):
10         plt_idx = i * num_classes + y + 1
11         plt.subplot(samples_per_class, num_classes, plt_idx)
12         plt.imshow(X_train[idx].astype('uint8'))
13         plt.axis('off')
14         if i == 0:
15             plt.title(cls)
16 plt.show()
```



In [4]:

```
1 # Split the data into train, val, and test sets. In addition we will
2 # create a small development set as a subset of the training data;
3 # we can use this for development so our code runs faster.
4 num_training = 49000
5 num_validation = 1000
6 num_test = 1000
7 num_dev = 500
8
9 # Our validation set will be num_validation points from the original
10 # training set.
11 mask = range(num_training, num_training + num_validation)
12 X_val = X_train[mask]
13 y_val = y_train[mask]
14
15 # Our training set will be the first num_train points from the original
16 # training set.
17 mask = range(num_training)
18 X_train = X_train[mask]
19 y_train = y_train[mask]
20
21 # We will also make a development set, which is a small subset of
22 # the training set.
23 mask = np.random.choice(num_training, num_dev, replace=False)
24 X_dev = X_train[mask]
25 y_dev = y_train[mask]
26
27 # We use the first num_test points of the original test set as our
28 # test set.
29 mask = range(num_test)
30 X_test = X_test[mask]
31 y_test = y_test[mask]
32
33 print('Train data shape: ', X_train.shape)
34 print('Train labels shape: ', y_train.shape)
35 print('Validation data shape: ', X_val.shape)
36 print('Validation labels shape: ', y_val.shape)
37 print('Test data shape: ', X_test.shape)
38 print('Test labels shape: ', y_test.shape)
```

Train data shape: (49000, 32, 32, 3)

Train labels shape: (49000,)

Validation data shape: (1000, 32, 32, 3)

Validation labels shape: (1000,)

Test data shape: (1000, 32, 32, 3)

Test labels shape: (1000,)

In [5]:

```
1 # Preprocessing: reshape the image data into rows
2 X_train = np.reshape(X_train, (X_train.shape[0], -1))
3 X_val = np.reshape(X_val, (X_val.shape[0], -1))
4 X_test = np.reshape(X_test, (X_test.shape[0], -1))
5 X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))
6
7 # As a sanity check, print out the shapes of the data
8 print('Training data shape: ', X_train.shape)
9 print('Validation data shape: ', X_val.shape)
10 print('Test data shape: ', X_test.shape)
11 print('dev data shape: ', X_dev.shape)
```

Training data shape: (49000, 3072)

Validation data shape: (1000, 3072)


Test data shape: (1000, 3072)

dev data shape: (500, 3072)

In [6]:

```
1 # Preprocessing: subtract the mean image
2 # first: compute the image mean based on the training data
3 mean_image = np.mean(X_train, axis=0)
4 print(mean_image[:10]) # print a few of the elements
5 plt.figure(figsize=(4,4))
6 plt.imshow(mean_image.reshape((32,32,3)).astype('uint8')) # visualize the mean image
7 plt.show()
```

```
[130.64189796 135.98173469 132.47391837 130.05569388 135.34804082
 131.75402041 130.96055102 136.14328571 132.47636735 131.48467347]
```



In [7]:

```
1 # second: subtract the mean image from train and test data
2 X_train -= mean_image
3 X_val -= mean_image
4 X_test -= mean_image
5 X_dev -= mean_image
```

In [8]:

```
1 # third: append the bias dimension of ones (i.e. bias trick) so that our SVM
2 # only has to worry about optimizing a single weight matrix W.
3 X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
4 X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
5 X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
6 X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])
7
8 print(X_train.shape, X_val.shape, X_test.shape, X_dev.shape)
```

(49000, 3073) (1000, 3073) (1000, 3073) (500, 3073)

SVM Classifier

Your code for this section will all be written inside **cs175/classifiers/linear_svm.py**.

As you can see, we have prefilled the function `compute_loss_naive` which uses for loops to evaluate the multiclass SVM loss function.

In [20]:

```
1 # Evaluate the naive implementation of the loss we provided for you:
2 from cs175.classifiers.linear_svm import svm_loss_naive
3 import time
4
5 # generate a random SVM weight matrix of small numbers
6 W = np.random.randn(3073, 10) * 0.0001
7
8 loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.000005)
9 print('loss: %f' % (loss, ))
```

loss: 8.818681

The `grad` returned from the function above is right now all zero. Derive and implement the gradient for the SVM cost function and implement it inline inside the function `svm_loss_naive`. You will find it helpful to interleave your new code inside the existing function.

To check that you have correctly implemented the gradient correctly, you can numerically estimate the gradient of the loss function and compare the numeric estimate to the gradient that you computed. We have provided code that does this for you:

In [10]:

```
1 # Once you've implemented the gradient, recompute it with the code below
2 # and gradient check it with the function we provided for you
3
4 # Compute the loss and its gradient at W.
5 loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.0)
6
7 # Numerically compute the gradient along several randomly chosen dimensions, and
8 # compare them with your analytically computed gradient. The numbers should match
9 # almost exactly along all dimensions.
10 from cs175.gradient_check import grad_check_sparse
11 f = lambda w: svm_loss_naive(w, X_dev, y_dev, 0.0)[0]
12 grad_numerical = grad_check_sparse(f, W, grad)
13
14 # do the gradient check once again with regularization turned on
15 # you didn't forget the regularization gradient did you?
16 loss, grad = svm_loss_naive(W, X_dev, y_dev, 5e1)
17 f = lambda w: svm_loss_naive(w, X_dev, y_dev, 5e1)[0]
18 grad_numerical = grad_check_sparse(f, W, grad)
```

```
numerical: 2.288954 analytic: 2.288954, relative error: 1.606896e-10
numerical: 14.059200 analytic: 14.059200, relative error: 1.502806e-11
numerical: 15.106084 analytic: 15.106084, relative error: 2.052577e-11
numerical: 27.875631 analytic: 27.875631, relative error: 5.368778e-12
numerical: -6.852777 analytic: -6.852777, relative error: 8.841052e-11
numerical: -9.140862 analytic: -9.140862, relative error: 6.045457e-12
numerical: -2.403386 analytic: -2.325934, relative error: 1.637695e-02
numerical: -21.631042 analytic: -21.631042, relative error: 2.008336e-11
numerical: 10.217683 analytic: 10.235062, relative error: 8.497107e-04
numerical: -9.348563 analytic: -9.348563, relative error: 4.654219e-11
numerical: -14.409265 analytic: -14.423620, relative error: 4.978610e-04
numerical: -6.955932 analytic: -6.961004, relative error: 3.644590e-04
numerical: 14.667952 analytic: 14.678718, relative error: 3.668516e-04
numerical: 3.654761 analytic: 3.670569, relative error: 2.157927e-03
numerical: -1.473110 analytic: -1.476481, relative error: 1.142641e-03
numerical: -13.633525 analytic: -13.666144, relative error: 1.194845e-03
numerical: 4.271784 analytic: 4.270527, relative error: 1.471787e-04
numerical: -14.681536 analytic: -14.682217, relative error: 2.317809e-05
numerical: 2.534981 analytic: 2.534381, relative error: 1.183488e-04
numerical: -24.472695 analytic: -24.461725, relative error: 2.241899e-04
```

Inline Question 1:

It is possible that once in a while a dimension in the gradcheck will not match exactly. What could such a discrepancy be caused by? Is it a reason for concern? What is a simple example in one dimension where a gradient check could fail? *Hint: the SVM loss function is not strictly speaking differentiable*

Your Answer: the gradient check would find at the turning point the value can not be differentiable. It is a reason for concern, one example is when it finds central minimum but not global minimum.

In [11]:

```
1 # Next implement the function svm_loss_vectorized; for now only compute the loss;
2 # we will implement the gradient in a moment.
3 tic = time.time()
4 loss_naive, grad_naive = svm_loss_naive(W, X_dev, y_dev, 0.000005)
5 toc = time.time()
6 print('Naive loss: %e computed in %fs' % (loss_naive, toc - tic))
7
8 from cs175.classifiers.linear_svm import svm_loss_vectorized
9 tic = time.time()
10 loss_vectorized, _ = svm_loss_vectorized(W, X_dev, y_dev, 0.000005)
11 toc = time.time()
12 print('Vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))
13
14 # The losses should match but your vectorized implementation should be much faster.
15 print('difference: %f' % (loss_naive - loss_vectorized))
```

Naive loss: 9.526596e+00 computed in 0.061021s
Vectorized loss: 9.526596e+00 computed in 0.003000s
difference: 0.000000

In [12]:

```
1 # Complete the implementation of svm_loss_vectorized, and compute the gradient
2 # of the loss function in a vectorized way.
3
4 # The naive implementation and the vectorized implementation should match, but
5 # the vectorized version should still be much faster.
6 tic = time.time()
7 _, grad_naive = svm_loss_naive(W, X_dev, y_dev, 0.000005)
8 toc = time.time()
9 print('Naive loss and gradient: computed in %fs' % (toc - tic))
10
11 tic = time.time()
12 _, grad_vectorized = svm_loss_vectorized(W, X_dev, y_dev, 0.000005)
13 toc = time.time()
14 print('Vectorized loss and gradient: computed in %fs' % (toc - tic))
15
16 # The loss is a single number, so it is easy to compare the values computed
17 # by the two implementations. The gradient on the other hand is a matrix, so
18 # we use the Frobenius norm to compare them.
19 difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
20 print('difference: %f' % difference)
```

Naive loss and gradient: computed in 0.072016s
Vectorized loss and gradient: computed in 0.002001s
difference: 0.000000

Stochastic Gradient Descent

We now have vectorized and efficient expressions for the loss, the gradient and our gradient matches the numerical gradient. We are therefore ready to do SGD to minimize the loss.


In [23]:

```
1 # In the file linear_classifier.py, implement SGD in the function
2 # LinearClassifier.train() and then run it with the code below.
3 from cs175.classifiers import LinearSVM
4 svm = LinearSVM()
5 tic = time.time()
6 loss_hist = svm.train(X_train, y_train, learning_rate=1e-7, reg=2.5e4,
7                       num_iters=1500, verbose=True)
8 toc = time.time()
9 print('That took %fs' % (toc - tic))
```

```
iteration 0 / 1500: loss 779.678451
iteration 100 / 1500: loss 468.502127
iteration 200 / 1500: loss 283.802481
iteration 300 / 1500: loss 173.096316
iteration 400 / 1500: loss 106.460575
iteration 500 / 1500: loss 65.548422
iteration 600 / 1500: loss 42.409072
iteration 700 / 1500: loss 27.415613
iteration 800 / 1500: loss 18.447036
iteration 900 / 1500: loss 13.005289
iteration 1000 / 1500: loss 9.915129
iteration 1100 / 1500: loss 8.222058
iteration 1200 / 1500: loss 7.547878
iteration 1300 / 1500: loss 6.449511
iteration 1400 / 1500: loss 6.534033
That took 4.404796s
```


In [24]:

```
1 # A useful debugging strategy is to plot the loss as a function of
2 # iteration number:
3 plt.plot(loss_hist)
4 plt.xlabel('Iteration number')
5 plt.ylabel('Loss value')
6 plt.show()
```



In [25]:

```
1 # Write the LinearSVM.predict function and evaluate the performance on both the
2 # training and validation set
3 y_train_pred = svm.predict(X_train)
4 print('training accuracy: %f' % (np.mean(y_train == y_train_pred), ))
5 y_val_pred = svm.predict(X_val)
6 print('validation accuracy: %f' % (np.mean(y_val == y_val_pred), ))
```

training accuracy: 0.381265
validation accuracy: 0.373000

In [32]:


```
1 # Use the validation set to tune hyperparameters (regularization strength and
2 # learning rate). You should experiment with different ranges for the learning
3 # rates and regularization strengths; if you are careful you should be able to
4 # get a classification accuracy of about 0.4 on the validation set.
5 learning_rates = [1e-7, 2e-7, 3e-7, 4e-7]
6 regularization_strengths = [2e4, 3e4, 4e4, 5e4]
7
8 # results is dictionary mapping tuples of the form
9 # (learning_rate, regularization_strength) to tuples of the form
10 # (training_accuracy, validation_accuracy). The accuracy is simply the fraction
11 # of data points that are correctly classified.
12 results = {}
13 best_val = -1 # The highest validation accuracy that we have seen so far.
14 best_svm = None # The LinearSVM object that achieved the highest validation rate.
15
16 #####
17 # TODO: #
18 # Write code that chooses the best hyperparameters by tuning on the validation #
19 # set. For each combination of hyperparameters, train a linear SVM on the #
20 # training set, compute its accuracy on the training and validation sets, and #
21 # store these numbers in the results dictionary. In addition, store the best #
22 # validation accuracy in best_val and the LinearSVM object that achieves this #
23 # accuracy in best_svm. #
24 # #
25 # Hint: You should use a small value for num_iters as you develop your #
26 # validation code so that the SVMs don't take much time to train; once you are #
27 # confident that your validation code works, you should rerun the validation #
28 # code with a larger value for num_iters. #
29 #####
30
31 for i in learning_rates:
32     for j in regularization_strengths:
33         s1 = LinearSVM()
34
35         s1.train(X_train, y_train, learning_rate=i, reg=j, num_iters=1500)
36
37         ytr = s1.predict(X_train)
38         train_accuracy = np.mean(y_train == ytr)
39
40         yval = s1.predict(X_val)
41         val_accuracy = np.mean(y_val == yval)
42
43         results[(i, j)] = (train_accuracy, val_accuracy)
44
45         if best_val < val_accuracy:
46             best_val = val_accuracy
47             best_svm = s1
48
49 #####
50 # #
51 # END OF YOUR CODE #
52 #####
53
54 # Print out results.
55 for lr, reg in sorted(results):
56     train_accuracy, val_accuracy = results[(lr, reg)]
57     print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
58         lr, reg, train_accuracy, val_accuracy))
59
```

```
60 | print('best validation accuracy achieved during cross-validation: %f' % best_val)
```

```
lr 1.000000e-07 reg 2.000000e+04 train accuracy: 0.376939 val accuracy: 0.382000
lr 1.000000e-07 reg 3.000000e+04 train accuracy: 0.377143 val accuracy: 0.385000
lr 1.000000e-07 reg 4.000000e+04 train accuracy: 0.369388 val accuracy: 0.369000
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.369286 val accuracy: 0.376000
lr 2.000000e-07 reg 2.000000e+04 train accuracy: 0.375082 val accuracy: 0.378000
lr 2.000000e-07 reg 3.000000e+04 train accuracy: 0.371327 val accuracy: 0.372000
lr 2.000000e-07 reg 4.000000e+04 train accuracy: 0.368245 val accuracy: 0.369000
lr 2.000000e-07 reg 5.000000e+04 train accuracy: 0.366143 val accuracy: 0.360000
lr 3.000000e-07 reg 2.000000e+04 train accuracy: 0.369898 val accuracy: 0.386000
lr 3.000000e-07 reg 3.000000e+04 train accuracy: 0.360082 val accuracy: 0.362000
lr 3.000000e-07 reg 4.000000e+04 train accuracy: 0.357143 val accuracy: 0.367000
lr 3.000000e-07 reg 5.000000e+04 train accuracy: 0.315082 val accuracy: 0.331000
lr 4.000000e-07 reg 2.000000e+04 train accuracy: 0.337469 val accuracy: 0.340000
lr 4.000000e-07 reg 3.000000e+04 train accuracy: 0.364694 val accuracy: 0.363000
lr 4.000000e-07 reg 4.000000e+04 train accuracy: 0.336122 val accuracy: 0.347000
lr 4.000000e-07 reg 5.000000e+04 train accuracy: 0.348020 val accuracy: 0.354000
best validation accuracy achieved during cross-validation: 0.386000
```

In [33]:

```
1 # Visualize the cross-validation results
2 import math
3 x_scatter = [math.log10(x[0]) for x in results]
4 y_scatter = [math.log10(x[1]) for x in results]
5
6 # plot training accuracy
7 marker_size = 100
8 colors = [results[x][0] for x in results]
9 plt.subplot(2, 1, 1)
10 plt.scatter(x_scatter, y_scatter, marker_size, c=colors)
11 plt.colorbar()
12 plt.xlabel('log learning rate')
13 plt.ylabel('log regularization strength')
14 plt.title('CIFAR-10 training accuracy')
15
16 # plot validation accuracy
17 colors = [results[x][1] for x in results] # default size of markers is 20
18 plt.subplot(2, 1, 2)
19 plt.scatter(x_scatter, y_scatter, marker_size, c=colors)
20 plt.colorbar()
21 plt.xlabel('log learning rate')
22 plt.ylabel('log regularization strength')
23 plt.title('CIFAR-10 validation accuracy')
24 plt.show()
```



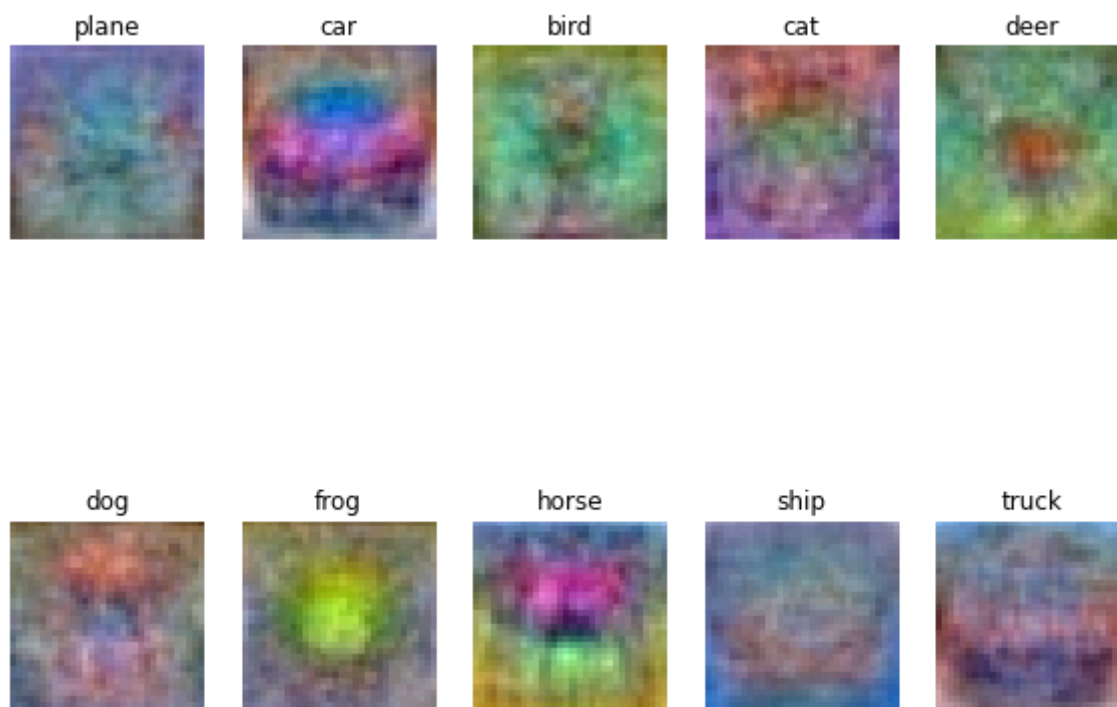
In [34]:

```
1 # Evaluate the best svm on test set
2 y_test_pred = best_svm.predict(X_test)
3 test_accuracy = np.mean(y_test == y_test_pred)
4 print('linear SVM on raw pixels final test set accuracy: %f' % test_accuracy)
```

linear SVM on raw pixels final test set accuracy: 0.370000

In [35]:

```
1 # Visualize the learned weights for each class.
2 # Depending on your choice of learning rate and regularization strength, these may
3 # or may not be nice to look at.
4 w = best_svm.W[:-1,:] # strip out the bias
5 w = w.reshape(32, 32, 3, 10)
6 w_min, w_max = np.min(w), np.max(w)
7 classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
8 for i in range(10):
9     plt.subplot(2, 5, i + 1)
10    # Rescale the weights to be between 0 and 255
11    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
12    plt.imshow(wimg.astype('uint8'))
13    plt.axis('off')
14    plt.title(classes[i])
```



Inline question 2:

Describe what your visualized SVM weights look like, and offer a brief explanation for why they look the way that they do.

Your answer: the visualized svm shows the relation between real images and trained svm. Car picture has more content in the middle and we can see the frame from visualized pictures.

